

In [1]:

```
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
import string
import math
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

import plotly
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
from collections import Counter
```

1.1 Reading Data

Splitting data into Train, Cross Validation and Test

In [2]:

```
preprocessed_data = pd.read_csv('preprocessed_data.csv')
preprocessed_data.head(2)
```

Out[2]:

	Unnamed: 0	Unnamed: 0.1	id	teacher_id	teacher
0	0	8393	p205479	2bf07ba08945e5d8b2a3f269b2b3cfe5	
1	1	37728	p043609	3f60494c61921b3b43ab61bdde2904df	

2 rows × 21 columns

In [3]:

```
y = preprocessed_data['project_is_approved'].values
X = preprocessed_data.drop(['project_is_approved'], axis=1)
X.shape
```

Out[3]:

(109248, 20)

In [4]:

```
# train test split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, stratify=y)
X_test.shape
```

Out[4]:

(36052, 20)

1.4 Encoding Categorical and Numerical features

1.4.1 encoding categorical features: clean_categories

In [5]:

```
feature_names=[]
feature_names_tfidf=[]

vectorizer = CountVectorizer()
vectorizer.fit(X_train['clean_categories'].values) # fit has to happen only on training data

# we use the fitted CountVectorizer to convert the text to vector
X_train_cc_oh = vectorizer.transform(X_train['clean_categories'].values)
#X_cv_cc_oh = vectorizer.transform(X_cv['clean_categories'].values)
X_test_cc_oh = vectorizer.transform(X_test['clean_categories'].values)

print("After vectorizations")
print(X_train_cc_oh.shape, y_train.shape)
#print(X_cv_cc_oh.shape, y_cv.shape)
print(X_test_cc_oh.shape, y_test.shape)
print(vectorizer.get_feature_names())

feature_names.extend(vectorizer.get_feature_names())
```

```
After vectorizations
(73196, 9) (73196,)
(36052, 9) (36052,)
['appliedlearning', 'care_hunger', 'health_sports', 'history_civics', 'literacy_language', 'math_science', 'music_arts', 'specialneeds', 'warmth']
```

1.4.2 encoding categorical features: clean_subcategories

In [6]:

```
vectorizer = CountVectorizer()
vectorizer.fit(X_train['clean_subcategories'].values) # fit has to happen on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_csc_ohe = vectorizer.transform(X_train['clean_subcategories'].values)
#X_cv_csc_ohe = vectorizer.transform(X_cv['clean_subcategories'].values)
X_test_csc_ohe = vectorizer.transform(X_test['clean_subcategories'].values)

print("After vectorizations")
print(X_train_csc_ohe.shape, y_train.shape)
#print(X_cv_csc_ohe.shape, y_cv.shape)
print(X_test_csc_ohe.shape, y_test.shape)
print(vectorizer.get_feature_names())

feature_names.extend(vectorizer.get_feature_names())
```

```
After vectorizations
(73196, 30) (73196,)
(36052, 30) (36052,)
['appliedsciences', 'care_hunger', 'charactereducation', 'civics_government', 'college_careerprep', 'communityservice', 'earlydevelopment', 'economics', 'environmentalscience', 'esl', 'extracurricular', 'financialliteracy', 'foreignlanguages', 'gym_fitness', 'health_lifescience', 'health_wellness', 'history_geography', 'literacy', 'literature_writing', 'mathematics', 'music', 'nutritioneducation', 'other', 'parentinvolvement', 'performingarts', 'socialsciences', 'specialneeds', 'teamsports', 'visualarts', 'warmth']
```

1.4.3 encoding categorical features: school_state

In [7]:

```
vectorizer = CountVectorizer()
vectorizer.fit(X_train['school_state'].values) # fit has to happen only on training data

# we use the fitted CountVectorizer to convert the text to vector
X_train_state_ohe = vectorizer.transform(X_train['school_state'].values)
#X_cv_state_ohe = vectorizer.transform(X_cv['school_state'].values)
X_test_state_ohe = vectorizer.transform(X_test['school_state'].values)

print("After vectorizations")
print(X_train_state_ohe.shape, y_train.shape)
#print(X_cv_state_ohe.shape, y_cv.shape)
print(X_test_state_ohe.shape, y_test.shape)
print(vectorizer.get_feature_names())

feature_names.extend(vectorizer.get_feature_names())
```

```
After vectorizations
(73196, 51) (73196,)
(36052, 51) (36052,)
['ak', 'al', 'ar', 'az', 'ca', 'co', 'ct', 'dc', 'de', 'fl',
'ga', 'hi', 'ia', 'id', 'il', 'in', 'ks', 'ky', 'la', 'ma', 'm
d', 'me', 'mi', 'mn', 'mo', 'ms', 'mt', 'nc', 'nd', 'ne', 'n
h', 'nj', 'nm', 'nv', 'ny', 'oh', 'ok', 'or', 'pa', 'ri', 's
c', 'sd', 'tn', 'tx', 'ut', 'va', 'vt', 'wa', 'wi', 'wv', 'w
y']
```

1.4.4 encoding categorical features: teacher_prefix

In [8]:

```
vectorizer = CountVectorizer()
vectorizer.fit(X_train['teacher_prefix'].values) # fit has to happen only on

# we use the fitted CountVectorizer to convert the text to vector
X_train_teacher_ohe = vectorizer.transform(X_train['teacher_prefix'].values)
#X_cv_teacher_ohe = vectorizer.transform(X_cv['teacher_prefix'].values)
X_test_teacher_ohe = vectorizer.transform(X_test['teacher_prefix'].values)

print("After vectorizations")
print(X_train_teacher_ohe.shape, y_train.shape)
#print(X_cv_teacher_ohe.shape, y_cv.shape)
print(X_test_teacher_ohe.shape, y_test.shape)
print(vectorizer.get_feature_names())

feature_names.extend(vectorizer.get_feature_names())
```

```
After vectorizations
(73196, 5) (73196,)
(36052, 5) (36052,)
['dr', 'mr', 'mrs', 'ms', 'teacher']
```

1.4.5 encoding categorical features: project_grade_category

In [9]:

```
vectorizer = CountVectorizer()
vectorizer.fit(X_train['project_grade_category'].values) # fit has to happen

# we use the fitted CountVectorizer to convert the text to vector
X_train_grade_ohe = vectorizer.transform(X_train['project_grade_category'].values)
#X_cv_grade_ohe = vectorizer.transform(X_cv['project_grade_category'].values)
X_test_grade_ohe = vectorizer.transform(X_test['project_grade_category'].values)

print("After vectorizations")
print(X_train_grade_ohe.shape, y_train.shape)
#print(X_cv_grade_ohe.shape, y_cv.shape)
print(X_test_grade_ohe.shape, y_test.shape)
print(vectorizer.get_feature_names())

feature_names.extend(vectorizer.get_feature_names())
```

```
After vectorizations
(73196, 4) (73196,)
(36052, 4) (36052,)
['grades_3_5', 'grades_6_8', 'grades_9_12', 'grades_prek_2']
```

1.4.6 encoding numerical features: price

In [10]:

```
from sklearn.preprocessing import Normalizer
normalizer = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
normalizer.fit(X_train['price'].values.reshape(1, -1))

X_train_price_norm = normalizer.transform(X_train['price'].values.reshape(1, -1))
#X_cv_price_norm = normalizer.transform(X_cv['price'].values.reshape(1, -1))
X_test_price_norm = normalizer.transform(X_test['price'].values.reshape(1, -1))

print("After vectorizations")
print(X_train_price_norm.shape, y_train.shape)
#print(X_cv_price_norm.shape, y_cv.shape)
print(X_test_price_norm.shape, y_test.shape)
print(X_train_price_norm)
print((X_test_price_norm))

feature_names.extend(['price'])
```

```
After vectorizations
(1, 73196) (73196,)
(1, 36052) (36052,)
[[0.00254418 0.00330246 0.00248577 ... 0.00180197 0.00020543
 0.00501878]]
[[0.00180586 0.00431706 0.00514938 ... 0.00215565 0.00299181
 0.00708843]]
```

1.4.7 encoding numerical features: teacher_number_of_previously_posted_projects

In [11]:

```
from sklearn.preprocessing import Normalizer
normalizer = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
normalizer.fit(X_train['teacher_number_of_previously_posted_projects'].values)

X_train_ppp_norm = normalizer.transform(X_train['teacher_number_of_previously_posted_projects'].values)
#X_cv_price_norm = normalizer.transform(X_cv['price'].values.reshape(1,-1))
X_test_ppp_norm = normalizer.transform(X_test['teacher_number_of_previously_posted_projects'].values)

print("After vectorizations")
print(X_train_ppp_norm.shape, y_train.shape)
#print(X_cv_price_norm.shape, y_cv.shape)
print(X_test_ppp_norm.shape, y_test.shape)

feature_names.extend(['teacher_number_of_previously_posted_projects'])
feature_names_tfidf.extend(feature_names)
```

```
After vectorizations
(1, 73196) (73196,)
(1, 36052) (36052,)
```

1.5 Vectorizing Text features

1.5.1 Vectorizing using BOW

Essay

In [12]:

```
print(X_train.shape, y_train.shape)
#print(X_cv.shape, y_cv.shape)
print(X_test.shape, y_test.shape)

print("\n\n")

vectorizer = CountVectorizer(min_df=10,ngram_range=(1,4), max_features=5000)
vectorizer.fit(X_train['essay'].values) # fit has to happen only on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_essay_bow = vectorizer.transform(X_train['essay'].values)
#X_cv_essay_bow = vectorizer.transform(X_cv['essay'].values)
X_test_essay_bow = vectorizer.transform(X_test['essay'].values)

print("After vectorizations")
print(X_train_essay_bow.shape, y_train.shape)
#print(X_cv_essay_bow.shape, y_cv.shape)
print(X_test_essay_bow.shape, y_test.shape)

feature_names.extend(vectorizer.get_feature_names())
```

```
(73196, 20) (73196,)
(36052, 20) (36052,)
```

```
After vectorizations
(73196, 5000) (73196,)
(36052, 5000) (36052,)
```

project_title

In [13]:

```
vectorizer = CountVectorizer(min_df=10,ngram_range=(1,4), max_features=5000)
vectorizer.fit(X_train['project_title'].values) # fit has to happen only on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_titles_bow = vectorizer.transform(X_train['project_title'].values)
#X_cv_titles_bow = vectorizer.transform(X_cv['project_title'].values)
X_test_titles_bow = vectorizer.transform(X_test['project_title'].values)

print("After vectorizations")
print(X_train_titles_bow.shape, y_train.shape)
#print(X_cv_titles_bow.shape, y_cv.shape)
print(X_test_titles_bow.shape, y_test.shape)

feature_names.extend(vectorizer.get_feature_names())
```

After vectorizations
(73196, 5000) (73196,)
(36052, 5000) (36052,)

project_resource_summary

In [14]:

```
vectorizer = CountVectorizer(min_df=10,ngram_range=(1,4), max_features=5000)
vectorizer.fit(X_train['project_resource_summary'].values) # fit has to happen only on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_psr_bow = vectorizer.transform(X_train['project_resource_summary'].values)
#X_cv_psr_bow = vectorizer.transform(X_cv['project_resource_summary'].values)
X_test_psr_bow = vectorizer.transform(X_test['project_resource_summary'].values)

print("After vectorizations")
print(X_train_psr_bow.shape, y_train.shape)
#print(X_cv_psr_bow.shape, y_cv.shape)
print(X_test_psr_bow.shape, y_test.shape)

feature_names.extend(vectorizer.get_feature_names())
```

After vectorizations
(73196, 5000) (73196,)
(36052, 5000) (36052,)

1.5.2 Vectorizing using TFIDF

essay

In [15]:

```
from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer(min_df=10,ngram_range=(1,4), max_features=5000)
vectorizer.fit(X_train['essay'].values)

X_train_essay_tfidf = vectorizer.transform(X_train['essay'].values)
#X_cv_essay_tfidf = vectorizer.transform(X_cv['essay'].values)
X_test_essay_tfidf = vectorizer.transform(X_test['essay'].values)

print("After vectorizations")
print(X_train_essay_tfidf.shape, y_train.shape)
#print(X_cv_essay_tfidf.shape, y_cv.shape)
print(X_test_essay_tfidf.shape, y_test.shape)

feature_names_tfidf.extend(vectorizer.get_feature_names())
```

After vectorizations
(73196, 5000) (73196,)
(36052, 5000) (36052,)

project_title

In [16]:

```
from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer(min_df=10,ngram_range=(1,4), max_features=5000)
vectorizer.fit(X_train['project_title'].values)

X_train_titles_tfidf = vectorizer.transform(X_train['project_title'].values)
#X_cv_titles_tfidf = vectorizer.transform(X_cv['project_title'].values)
X_test_titles_tfidf = vectorizer.transform(X_test['project_title'].values)

print("After vectorizations")
print(X_train_titles_tfidf.shape, y_train.shape)
#print(X_cv_titles_tfidf.shape, y_cv.shape)
print(X_test_titles_tfidf.shape, y_test.shape)

feature_names_tfidf.extend(vectorizer.get_feature_names())
```

After vectorizations
(73196, 5000) (73196,)
(36052, 5000) (36052,)

project_resource_summary

In [17]:

```
from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer(min_df=10,ngram_range=(1,4), max_features=5000)
vectorizer.fit(X_train['project_resource_summary'].values)

X_train_prs_tfidf = vectorizer.transform(X_train['project_resource_summary'].values)
#X_cv_prs_tfidf = vectorizer.transform(X_cv['project_resource_summary'].values)
X_test_prs_tfidf = vectorizer.transform(X_test['project_resource_summary'].values)

print("After vectorizations")
print(X_train_prs_tfidf.shape, y_train.shape)
#print(X_cv_prs_tfidf.shape, y_cv.shape)
print(X_test_prs_tfidf.shape, y_test.shape)

feature_names_tfidf.extend(vectorizer.get_feature_names())
```

```
After vectorizations
(73196, 5000) (73196,)
(36052, 5000) (36052,)
```

Merging all the categorical and numerical features with variations of text features

In [18]:

```
# merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039
from scipy.sparse import hstack

X_train_bow_matrix = hstack((X_train_cc_ohe, X_train_csc_ohe, X_train_state_ohe,
                             X_train_grade_ohe, X_train_price_norm.reshape(-1,1),
                             X_train_essay_bow, X_train_titles_bow, X_train_psr_bow))

X_test_bow_matrix = hstack((X_test_cc_ohe, X_test_csc_ohe, X_test_state_ohe,
                             X_test_price_norm.reshape(-1,1), X_test_ppp_norm,
                             X_test_essay_bow, X_test_titles_bow,
                             X_test_psr_bow)).tocsr()

print("Final Data matrix")
print(X_train_bow_matrix.shape, y_train.shape)
#print(X_cv_bow_matrix.shape, y_cv.shape)
print(X_test_bow_matrix.shape, y_test.shape)

print(len(feature_names))
```

```
Final Data matrix
(73196, 15101) (73196,)
(36052, 15101) (36052,)
15101
```

In [19]:

```
# merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039
from scipy.sparse import hstack

X_train_tfidf_matrix = hstack((X_train_cc_ohe, X_train_csc_ohe, X_train_state_ohe,
                               X_train_price_norm.reshape(-1,1), X_train_ppp_norm,
                               X_train_titles_tfidf, X_train_essay_tfidf, X_train_review_tfidf))

X_test_tfidf_matrix = hstack((X_test_cc_ohe, X_test_csc_ohe, X_test_state_ohe,
                              X_test_price_norm.reshape(-1,1), X_test_ppp_norm,
                              X_test_titles_tfidf, X_test_prs_tfidf)).tocsr()

print("Final Data matrix")
print(X_train_tfidf_matrix.shape, y_train.shape)
#print(X_cv_tfidf_matrix.shape, y_cv.shape)
print(X_test_tfidf_matrix.shape, y_test.shape)

print(len(feature_names_tfidf))
```

```
Final Data matrix
(73196, 15101) (73196,)
(36052, 15101) (36052,)
15101
```

Finding Best Hyper parameter using K-Fold CV on BOW representation of text features

In [20]:

```
# https://scikit-learn.org/stable/modules/generated/sklearn.model\_selection.C
from sklearn.model_selection import GridSearchCV
from scipy.stats import randint as sp_randint
from scipy.stats import norm as nm
from sklearn.model_selection import RandomizedSearchCV

import matplotlib.pyplot as plt
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import roc_auc_score

nb = MultinomialNB(class_prior = [0.5,0.5])
parameters = {'alpha':[0.001,0.01,0.1,1,10,10**2,10**3]}
clf1 = GridSearchCV(nb, parameters, cv=10, scoring='roc_auc', return_train_score=False)
clf1.fit(X_train_bow_matrix, y_train)

results = pd.DataFrame.from_dict(clf1.cv_results_)
results = results.sort_values(['param_alpha'])


train_auc=results['mean_train_score']
train_auc_std= results['std_train_score']
cv_auc = results['mean_test_score']
cv_auc_std= results['std_test_score']
K = results['param_alpha']

plt.plot(K, train_auc, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
# plt.gca().fill_between(K, train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2)

plt.plot(K, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
# plt.gca().fill_between(K, cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2)

plt.scatter(K, train_auc, label='Train AUC points')
plt.scatter(K, cv_auc, label='CV AUC points')

plt.legend()
plt.xlabel("K: hyperparameter")
plt.xscale('log')
plt.ylabel("AUC")
plt.title("Hyper parameter Vs AUC plot")
plt.grid()
plt.show()
results
```



Out[20]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha
0	0.229835	0.009303	0.013374	0.001419	0.001
1	0.227582	0.002220	0.013407	0.001494	0.01
2	0.229150	0.006135	0.013963	0.001355	0.1
3	0.226774	0.002843	0.013261	0.001734	1
4	0.229773	0.004796	0.013775	0.001595	10
5	0.223946	0.002794	0.013564	0.001286	100
6	0.228243	0.002483	0.012644	0.000888	1000

7 rows × 31 columns

Finding Best Hyper parameter using K-Fold CV on TFIDF representation of text features

In [21]:

```
# https://scikit-learn.org/stable/modules/generated/sklearn.model\_selection.C
from sklearn.model_selection import GridSearchCV
from scipy.stats import randint as sp_randint
from scipy.stats import norm as nm
from sklearn.model_selection import RandomizedSearchCV
import matplotlib.pyplot as plt
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import roc_auc_score

nb = MultinomialNB(class_prior = [0.5,0.5])
parameters = {'alpha':[0.001,0.01,0.1,1,10,10**2,10**3]}
clf = GridSearchCV(nb, parameters, cv=10, scoring='roc_auc', return_train_score=False)
clf.fit(X_train_tfidf_matrix, y_train)

results = pd.DataFrame.from_dict(clf.cv_results_)
results = results.sort_values(['param_alpha'])


train_auc= results['mean_train_score']
train_auc_std= results['std_train_score']
cv_auc = results['mean_test_score']
cv_auc_std= results['std_test_score']
K = results['param_alpha']

plt.plot(K, train_auc, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
# plt.gca().fill_between(K, train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2)

plt.plot(K, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
# plt.gca().fill_between(K, cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2)

plt.scatter(K, train_auc, label='Train AUC points')
plt.scatter(K, cv_auc, label='CV AUC points')

plt.legend()
plt.xlabel("K: hyperparameter")
plt.xscale('log')
plt.ylabel("AUC")
plt.title("Hyper parameter Vs AUC plot")
plt.grid()
plt.show()
results
```



Out[21]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha
0	0.232456	0.005960	0.014369	0.002146	0.001
1	0.231277	0.005977	0.013364	0.001194	0.01
2	0.229661	0.001024	0.012962	0.001412	0.1
3	0.229203	0.003328	0.013555	0.001335	1
4	0.231842	0.006450	0.013364	0.000917	10
5	0.231653	0.004571	0.013157	0.001388	100
6	0.229119	0.001947	0.013222	0.000931	1000

7 rows × 31 columns

In [22]:

```
def batch_predict(clf, data):
    # roc_auc_score(y_true, y_score) the 2nd parameter should be probability
    # not the predicted outputs

    y_data_pred = []
    tr_loop = data.shape[0] - data.shape[0]%1000
    # consider you X_tr shape is 49041, then your tr_loop will be 49041 - 490
    # in this for loop we will iterate until the last 1000 multiplier
    for i in range(0, tr_loop, 1000):
        y_data_pred.extend(clf.predict_proba(data[i:i+1000])[:,1])
    # we will be predicting for the last data points
    if data.shape[0]%1000 != 0:
        y_data_pred.extend(clf.predict_proba(data[tr_loop:])[:,1])

    return y_data_pred
```

In [23]:

```
# we are writing our own function for predict, with defined threshold
# we will pick a threshold that will give the least fpr
def find_best_threshold(threshold, fpr, tpr):
    t = threshold[np.argmax(tpr*(1-fpr))]
    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high
    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold", t)
    return t

def predict_with_best_t(proba, threshold):
    predictions = []
    for i in proba:
        if i >= threshold:
            predictions.append(1)
        else:
            predictions.append(0)
    return predictions
```

Applying NB with obtained best Alpha (Hyper parameter) on BOW

In [24]:


```
# https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc\_curve
from sklearn.metrics import roc_curve, auc

clf1 = MultinomialNB(alpha = 1, class_prior=[0.5,0.5])
clf1.fit(X_train_bow_matrix, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates
# not the predicted outputs

y_train_pred = batch_predict(clf1, X_train_bow_matrix)
y_test_pred = batch_predict(clf1, X_test_bow_matrix)

train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("TPR")
plt.ylabel("FPR")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()
```



Confusion Matrix with predicted and original labels for BOW

In [25]:

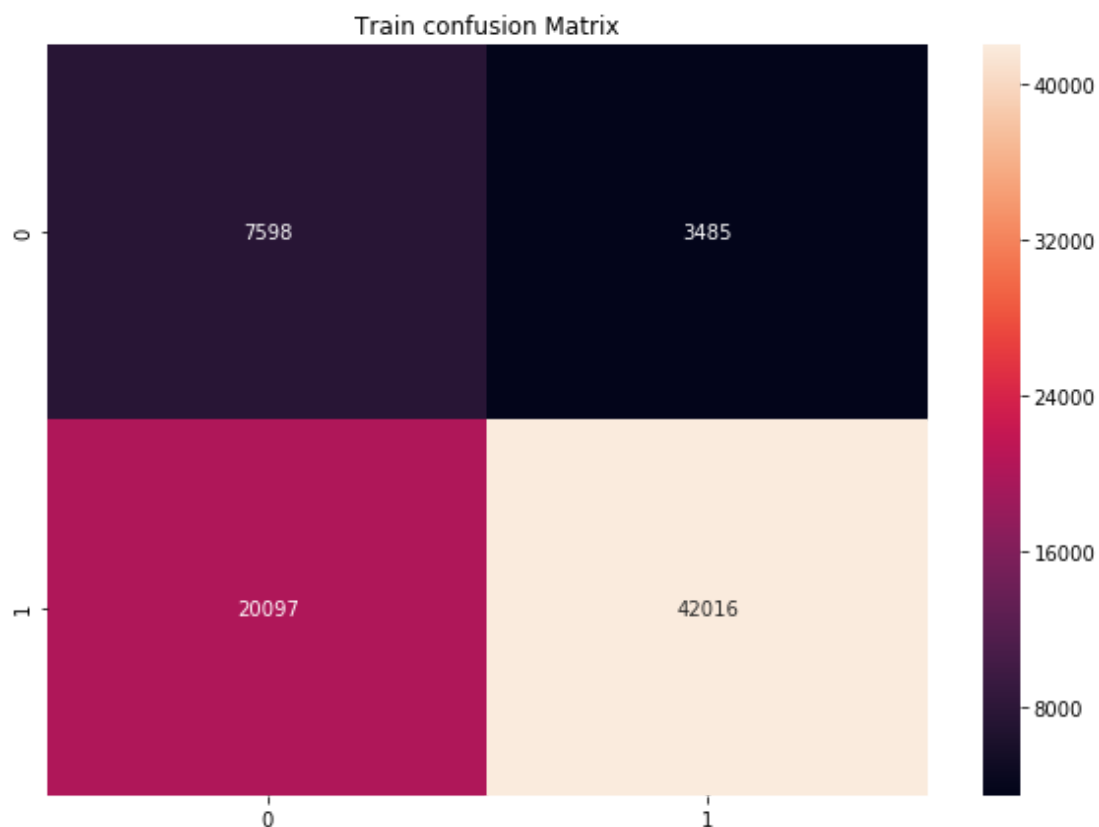
```
from sklearn.metrics import confusion_matrix
best_t = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
train = confusion_matrix(y_train, predict_with_best_t(y_train_pred, best_t))
test = confusion_matrix(y_test, predict_with_best_t(y_test_pred, best_t))


#https://stackoverflow.com/a/35572247

df_cm = pd.DataFrame(train, index = [i for i in range(2)], columns = [i for i in range(2)])
plt.figure(figsize = (10,7))
plt.title('Train confusion Matrix')
sns.heatmap(train, annot=True, fmt="d")
plt.show()

df_cm = pd.DataFrame(test, index = [i for i in range(2)], columns = [i for i in range(2)])
plt.figure(figsize = (10,7))
plt.title('Test confusion Matrix')
sns.heatmap(test, annot=True, fmt="d")
plt.show()
```

the maximum value of $tpr \cdot (1 - fpr)$ 0.46373956961336765 for threshold 0.448





Applying NB with obtained best Alpha (Hyper parameter) on TFIDF

In [26]:


```
# https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc\_curve
from sklearn.metrics import roc_curve, auc

clf = MultinomialNB(alpha = 1, class_prior=[0.5,0.5])
clf.fit(X_train_tfidf_matrix, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates
# not the predicted outputs

y_train_pred = batch_predict(clf, X_train_tfidf_matrix)
y_test_pred = batch_predict(clf, X_test_tfidf_matrix)

train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("TPR")
plt.ylabel("FPR")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()
```



Confusion Matrix with predicted and original labels for TFIDF

In [27]:

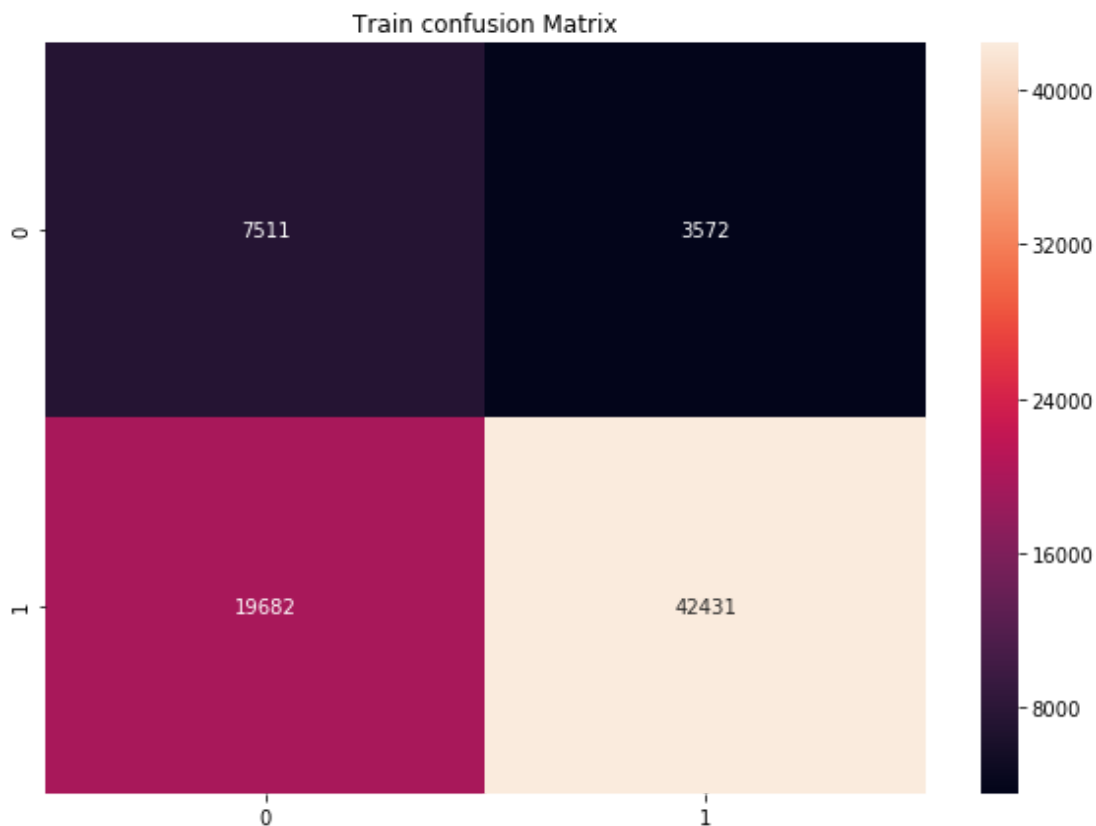
```
from sklearn.metrics import confusion_matrix
best_t = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
train = confusion_matrix(y_train, predict_with_best_t(y_train_pred, best_t))
test = confusion_matrix(y_test, predict_with_best_t(y_test_pred, best_t))

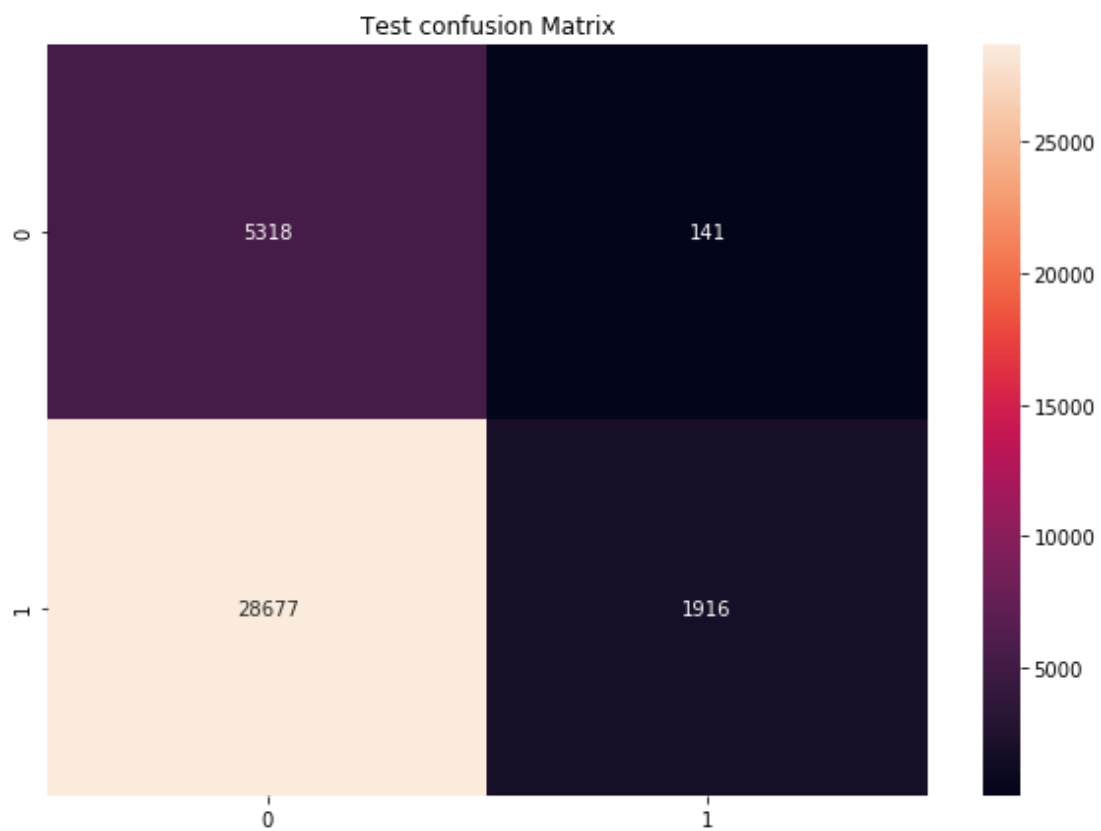
#https://stackoverflow.com/a/35572247

df_cm = pd.DataFrame(train, index = [i for i in range(2)], columns = [i for i in range(2)])
plt.figure(figsize = (10,7))
plt.title('Train confusion Matrix')
sns.heatmap(train, annot=True, fmt="d")
plt.show()

df_cm = pd.DataFrame(test, index = [i for i in range(2)], columns = [i for i in range(2)])
plt.figure(figsize = (10,7))
plt.title('Test confusion Matrix')
sns.heatmap(test, annot=True, fmt="d")
plt.show()
```

the maximum value of $tpr \cdot (1 - fpr)$ 0.4629575703867252 for threshold 0.501





Top 20 BOW features

In [28]:

```
max_ind_neg=np.argsort((clf1.feature_log_prob_)[0][::-1])[0:20]
top_neg=np.take(feature_names,max_ind_neg)
print('Number of features are : {}'.format(len(feature_names)))
print('\nIndices of top features to determine negative class label are :\n{')
print('\nNames of top features to determine negative class label are :\n{').format(

max_ind_pos=np.argsort((clf1.feature_log_prob_)[1][::-1])[0:20]
top_pos=np.take(feature_names,max_ind_pos)
print('\nIndices of top features to determine positive class label are :\n{')
print('\nNames of top features to determine positive class label are :\n{').format(
```

Number of features are : 15101

Indices of top features to determine negative class label are

```
:
[8918 9726 5764 6876 7482 8232 8231 8230 8216 8212 5712 5698 5
697 5696
 5691 9376 8002 9375 9374 9368]
```

Names of top features to determine negative class label are :

```
['space for' 'voices' 'can see' 'going green' 'language learne
rs'
 'our fingertips' 'our ears' 'our day' 'osmos' 'organizing'
 'building classroom' 'buddies' 'bronx' 'broken' 'bringing'
 'the world with' 'necessities' 'the world to' 'the world thro
ugh'
 'the win']
```

Indices of top features to determine positive class label are

```
:
[15010 7585 6364 8632 8716 9774 8856 6598 5054 8599
8280 6939
 6372 8674 6252 6889 9796 9671 6299 6282]
```

Names of top features to determine positive class label are :

```
['with hands on' 'let find' 'expanding' 'rescue' 'school suppl
ies for'
 'we like' 'signing day' 'for everything' 'would make'
 'reading writing and' 'out of' 'growing' 'exploration' 'rolli
ng'
 'electricity' 'google classroom' 'we ve' 'us be' 'engineering
in'
 'engage students']
```

Top 20 TFIDF features

In [29]:

```
max_ind_neg=np.argsort((clf.feature_log_prob_)[0][::-1])[0:20]
top_neg=np.take(feature_names_tfidf,max_ind_neg)
print('Number of features are : {}'.format(len(feature_names)))
print('\nIndices of top features to determine negative class label are :\n{')
print('\nNames of top features to determine negative class label are :\n{').format

max_ind_pos=np.argsort((clf.feature_log_prob_)[1][::-1])[0:20]
top_pos=np.take(feature_names_tfidf,max_ind_pos)
print('\nIndices of top features to determine positive class label are :\n{')
print('\nNames of top features to determine positive class label are :\n{').format
```

Number of features are : 15101

Indices of top features to determine negative class label are

:

[10838 11641 11640 11615 11613 11611 11602 14826 11600 11593 1
1588 14831

11517 11470 11464 11413 11412 11388 11347 11339]

Names of top features to determine negative class label are :

['carpet' 'growing' 'grow their' 'grades' 'grade level' 'googl
e classroom'

'glue and' 'use technology' 'gloves' 'give' 'get their wiggles'
'use to'

'for reading and' 'focusing' 'focus on their' 'fill' 'file'
'extra'

'ever' 'erasers and']

Indices of top features to determine positive class label are

:

[12585 13632 10596 10595 12182 10782 13599 11364 10299 10899 1
1359 10692

10054 10762 14866 13141 12844 12845 13139 10597]

Names of top features to determine positive class label are :

['need binders' 'stay focused and' 'balls for' 'balls and' 'ma
cbook air'

'build and' 'sports' 'experience in' 'and bouncy' 'charging'
'exercise'

'bins' 'yoga for' 'bouncy bands and' 'voice' 'place' 'need va
riety of'

'need variety of books' 'pillows' 'balls to']

Conclusion

In [30]:

```
# http://zetcode.com/python/prettytable/
from prettytable import PrettyTable

table = PrettyTable()
table.field_names = ["Vectorizer", "Model", "Hyper Parameter", "AUC"]

table.add_row(['BOW', 'Naive Bayes', 1, 0.7070])
table.add_row(['TFIDF', 'Naive Bayes', 1, 0.6509])
print(table)
```

Vectorizer	Model	Hyper Parameter	AUC
BOW	Naive Bayes	1	0.707
TFIDF	Naive Bayes	1	0.6509

Summary

- BOW vectorizer gave AUC 0.7070 with the best hyper parameter 1
- TFIDF vectorizer gave AUC 0.6509 with the best hyper parameter 1
- BOW vectorizer has better AUC compared to TFIDF vectorizer